
Note: You will be given time to solve the exercises during the tutorial.

Exercise 1.1: Estimate π with Monte Carlo

- a) Estimate π by “shooting” (i.e., drawing random numbers) N times uniformly on a square and counting the number of points hitting a disc target: the ratio of hits to N should correspond to the ratio of the areas of the target to the area you shoot on.
- b) Estimate the variance of the error for given N by repeating this a few times.
- c) Plot the variance of the error versus N on a log-log scale. What is the scaling of the error?
- d) Generalize your code to estimate the volume of a d -dimensional sphere. Does it run much slower for large d ?

Exercise 1.2: Importance sampling with Monte Carlo

The goal of this exercise is to find a Monte Carlo estimate of $I = \int_{a=0}^{\infty} dx \frac{e^{-x}}{1+(x-1)^2}$. (It is fine to cut off the upper limit at some value e.g. $b=10$.)

- a) Write a function that calculates I by using that $I = \bar{f} \times (b - a)$ with $f(x) = \frac{e^{-x}}{1+(x-1)^2}$. (You can find an estimate of \bar{f} by uniformly generating some $x_i \in [a, b]$ and averaging over $f(x_i)$.) Get an idea of the error of this estimate.
- b) Now we will use importance sampling. The p.d.f. $g(x) = \alpha e^{-\alpha x}$ turns out to be a good choice. In order to maximally improve our estimate of I , we will need the optimal value of α . Write a function that determines the optimal α , or use the function in the script `integral.py`.
- c) Use importance sampling to estimate I over the same number of samples as in (a). How much improved your result?

Exercise 1.3: Metropolis algorithm for the 2D Ising model

Download the script `metropolis.py` from the homepage, which implements the Metropolis algorithm for the classical 2D Ising model $H = -J \sum_{\langle i,j \rangle} \sigma_i \sigma_j$ with $J \equiv 1$. The 2D Ising model has a critical point at $T_c = 2J / \ln(1 + \sqrt{2}) \approx 2.269$.

- a) What is the script plotting?
- b) What are “typical” configurations at temperatures $T \gg T_c$, $T \approx T_c$ and $T \ll T_c$?

- c) Plot the energy E and specific heat C_V versus temperature T for different system sizes L .
- d) Adjust the script to measure the magnetization $M = \frac{1}{L^2} \sum_{i,j} \sigma_{i,j}$. Plot how M changes with simulation time (=the number of updates performed) for $T > T_c$, $T \approx T_c$ and $T < T_c$. Which time scales can you recognize? In which cases do you still get the correct expectation value $\langle M \rangle = 0$? Plot $\langle |M| \rangle$ (i.e. taking the absolute value of M before averaging) versus T to see the transition.
- e) Include a magnetic field h coupling to the spins with a term $H' = -h \sum_i \sigma_i$. Plot $\langle M \rangle$ versus T .

Bonus Some further ideas for playing around:

- Instead of restarting from a random state for each new β , re-use the last state of the previous simulation. You should still perform sweeps without measurements for the thermalization. Is it better to start with large β or small β ?
- Change the lattice.
- Optimize the code.
- ...